

AUTOMATED INCIDENT RESPONSE WITH OSQUERY AND LOKI

GEORGE ADAMS IV & ED WELCH

USE CASE

DETECTING AND ALERTING SSH CONNECTIONS.

2/38

This is the end goal we'll build towards as we learn a little more along the way.

FOLLOW ALONG

[HTTPS://GITHUB.COM/GEOWA4/LEARN-LOKI](https://github.com/GEOWA4/LEARN-LOKI)

OSQUERY

RELEASED BY FACEBOOK IN 2014

```
commit 73a32b729403b2f5a7c204b0f7cfb86fdfdd0a85
Author: mike@arpaia.co <mike@arpaia.co>
Date:   Wed Jul 30 17:35:19 2014 -0700
```

Initial commit

OSQUERY

WORKS ON MY MACHINE



OSQUERY

SQL INTERFACE TO YOUR ENDPOINTS

.schema processes

```
CREATE TABLE processes(  
  `pid` BIGINT, `name` TEXT, `path` TEXT, `cmdline` TEXT,  
  `disk_bytes_read` BIGINT, `disk_bytes_written` BIGINT,  
  ...  
  PRIMARY KEY (`pid`)  
) WITHOUT ROWID;
```

OSQUERY

```
select max(disk_bytes_read), pid, name, cmdline, cwd from processes;
```

max(disk_bytes_read)	pid	name	cmdline	cwd
10465280	5478	bash	-bash	/vagrant/demo

7/38

Who remembers how the
command to list what
processes are listening on
which ports?

OSQUERY

```
select p.name, l.port, l.protocol
from processes p inner join listening_ports l on p.pid = l.pid
where p.name = 'VBoxHeadless' and l.port <> 0;
```

name	port	protocol
VBoxHeadless	3000	6
VBoxHeadless	3100	6
VBoxHeadless	9080	6
VBoxHeadless	9090	6
VBoxHeadless	2222	6

OSQUERY

```
select c.name, p.port, p.host_port
from docker_containers c inner join docker_container_ports p
on c.id = p.id;
```

name	port	host_port
/demo_loki_1	80	0
/demo_loki_1	3100	3100
/demo_grafana_1	3000	3000
/demo_prometheus_1	9090	9090
/demo_promtail_1	9080	9080

OSQUERY

.schema last

```
CREATE TABLE last(  
  `username` TEXT, `time` INTEGER, `host` TEXT,  
  `pid` INTEGER, `tty` TEXT, `type` INTEGER  
);
```

OSQUERY

```
select * from last;
```

username	tty	pid	type	time	host
reboot	~	0	2	1566866107	4.15.0-52-generic
runlevel	~	53	1	1566866118	4.15.0-52-generic
	ttyS0	859	5	1566866119	
LOGIN	ttyS0	859	6	1566866119	
	tty1	879	5	1566866119	
LOGIN	tty1	879	6	1566866119	
vagrant	pts/0	5396	7	1566869034	10.0.2.2
vagrant	pts/1	6465	7	1566870878	10.0.2.2
	pts/1	6465	8	1566870880	
vagrant	pts/1	6571	7	1566870886	10.0.2.2
	pts/1	6571	8	1566870910	

OSQUERY – PACKS

```
{  
  "queries": {  
    "last": {  
      "query": "select * from last;",  
      "interval": "60",  
      "platform": "posix",  
      "version": "1.4.5",  
      "description": "..."  
    }  
  }  
}
```

OSQUERY – DECORATORS

```
{  
  "decorators": {  
    "load": [  
      "SELECT uuid AS host_uuid FROM system_info;",  
      "SELECT user AS username FROM logged_in_users ORDER BY time DESC LIMIT 1;"  
    ]  
  }  
}
```

13/38

Decorator results are added to pack results.

If you're on AWS, you can add a query for your EC2's tags.

OSQUERY – RESULTS

/var/log/osquery/osqueryd.results.log

```
{
  "name": "pack_incident-response_last",
  "hostIdentifier": "ubuntu-bionic",
  "calendarTime": "Tue Aug 27 01:55:13 2019 UTC",
  "decorations": {
    "host_uuid": "2401CCE9-23EA-4D4D-8C84-D5C8437EBE15",
    "username": "vagrant"
  },
  "columns": {
    "host": "10.0.2.2",
    "pid": "6465",
    "time": "1566870878",
    "tty": "pts/1",
    "type": "7",
    "username": "vagrant"
  },
  "action": "added"
}
```

OSQUERY – RECAP

IT'S BEEN AROUND A WHILE

CROSS-PLATFORM

IT'S JUST SQL

SCHEDULE QUERIES WITH 'PACKS'

LOKI

CLOUD-NATIVE LOG AGGREGATION

MADE BY GRAFANA

LOKI

PROMETHEUS-INSPIRED LOGGING FOR CLOUD NATIVES.

17/38

To know what this means, let's take a quick detour to Prometheus.

PROMETHEUS

METRICS COLLECTION VIA 'PULL'

TIME-SERIES DATA STORE

QUERYABLE VIA PROMQL

PROMETHEUS – PUSH VS. PULL

19/38

TODO: diagram push vs pull

PROMETHEUS – SCRAPING

```
scrape_configs:  
  - job_name: "promtail"  
    static_configs:  
      - targets:  
        - promtail:9080
```

PROMETHEUS – DATA STRUCTURE

METRICS HAVE LABELS IN ADDITION TO VALUES.

```
rss_enjoyment{track="tech", talk="osquery_loki"} 11
```

LOKI – DATA STRUCTURE

LOG ENTRIES HAVE LABELS, TOO.

```
{track="tech", talk="osquery_loki"} "best talk ever"  
{track="tech", talk="osquery_loki"} "i want to know more"  
{track="tech", talk="osquery_loki"} "i hear that one guy runs rocdev"
```

LOKI – LABELS

THE MATCHING LABELS ALLOW US TO SWITCH BACK AND FORTH FREELY.

LOKI – QUERY

LOGQL

```
$ logcli query --tail '{name="pack_incident-response_last"}'  
2019-08-29T03:01:37Z  
{filename="/var/log/osquery/osqueryd.results.log", job="osquery_results", name="pack_incident-response_last"}  
{  
  "name":"pack_incident-response_last",  
  "hostIdentifier":"ubuntu-bionic",  
  "calendarTime":"Thu Aug 29 03:01:37 2019 UTC",  
  "unixTime":1567047697,  
  "epoch":0,  
  "counter":115,  
  "decorations":{"host_uuid":"661449FD-E11A-462B-9EA9-63A3EE8F9BDC","username":"vagrant"},  
  "columns":{"host":"10.0.2.2","pid":"7404","time":"1567047680","tty":"pts/1","type":"7","username":"vagrant"},  
  "action":"added"  
}
```

24/38

This is why Loki is like Prometheus.
We read from them the same way
with the same labels.
But there's one more link with
Prometheus and Loki: metrics
extraction.

LOKI - COLLECTION

PROMTAIL

25/38

We've seen how Osquery generates logs.

We've seen how to read logs.
But how did they get there?

PROMTAIL

FORWARDS LOGS
AND
EXTRACTS METRICS

PROMTAIL – SCRAPING

```
clients:
  - url: http://loki:3100/api/prom/push
scrape_configs:
  - job_name: osquery
    static_configs:
      - targets:
          - localhost
    labels:
      job: osquery_results
      __path__: /var/log/osquery/osqueryd.results.log
```

27/38

Anything Osquery reports will be forwarded to Loki

PROMTAIL – RESULT REMINDER

```
{
  "name": "pack_incident-response_last",
  "hostIdentifier": "ubuntu-bionic",
  "calendarTime": "Thu Aug 29 03:01:37 2019 UTC",
  "unixTime": 1567047697,
  "epoch": 0,
  "counter": 115,
  "decorations": {
    "host_uuid": "661449FD-E11A-462B-9EA9-63A3EE8F9BDC",
    "username": "vagrant"
  },
  "columns": {
    "host": "10.0.2.2",
    "pid": "7404",
    "time": "1567047680",
    "tty": "pts/1",
    "type": "7",
    "username": "vagrant"
  },
  "action": "added"
}
```

PROMTAIL – PIPELINES

```
pipeline_stages:  
  - json:  
    expressions:  
      timestamp: unixTime  
      ip: columns.host  
      name: name  
      username: columns.username  
  - timestamp:  
    source: timestamp  
    format: Unix  
  - labels:  
    name: name
```

PROMTAIL - METRICS

```
pipeline_stages:
- ...
- metrics:
    last_logins:
        type: Counter
        description: count last logins
        source: name
        config:
            value: pack_incident-response_last
            action: inc
```

PROMTAIL – PROMETHEUS

```
scrape_configs:  
  - job_name: "promtail"  
    static_configs:  
      - targets:  
        - promtail:9080
```

WHERE WE ARE NOW

OSQUERY PRODUCING RESULTS

PROMTAIL FORWARDING TO LOKI

QUERY AND TAIL LOGS IN LOKI

PROMTAIL EXTRACTING METRICS

PROMETHEUS SCRAPING PROMTAIL

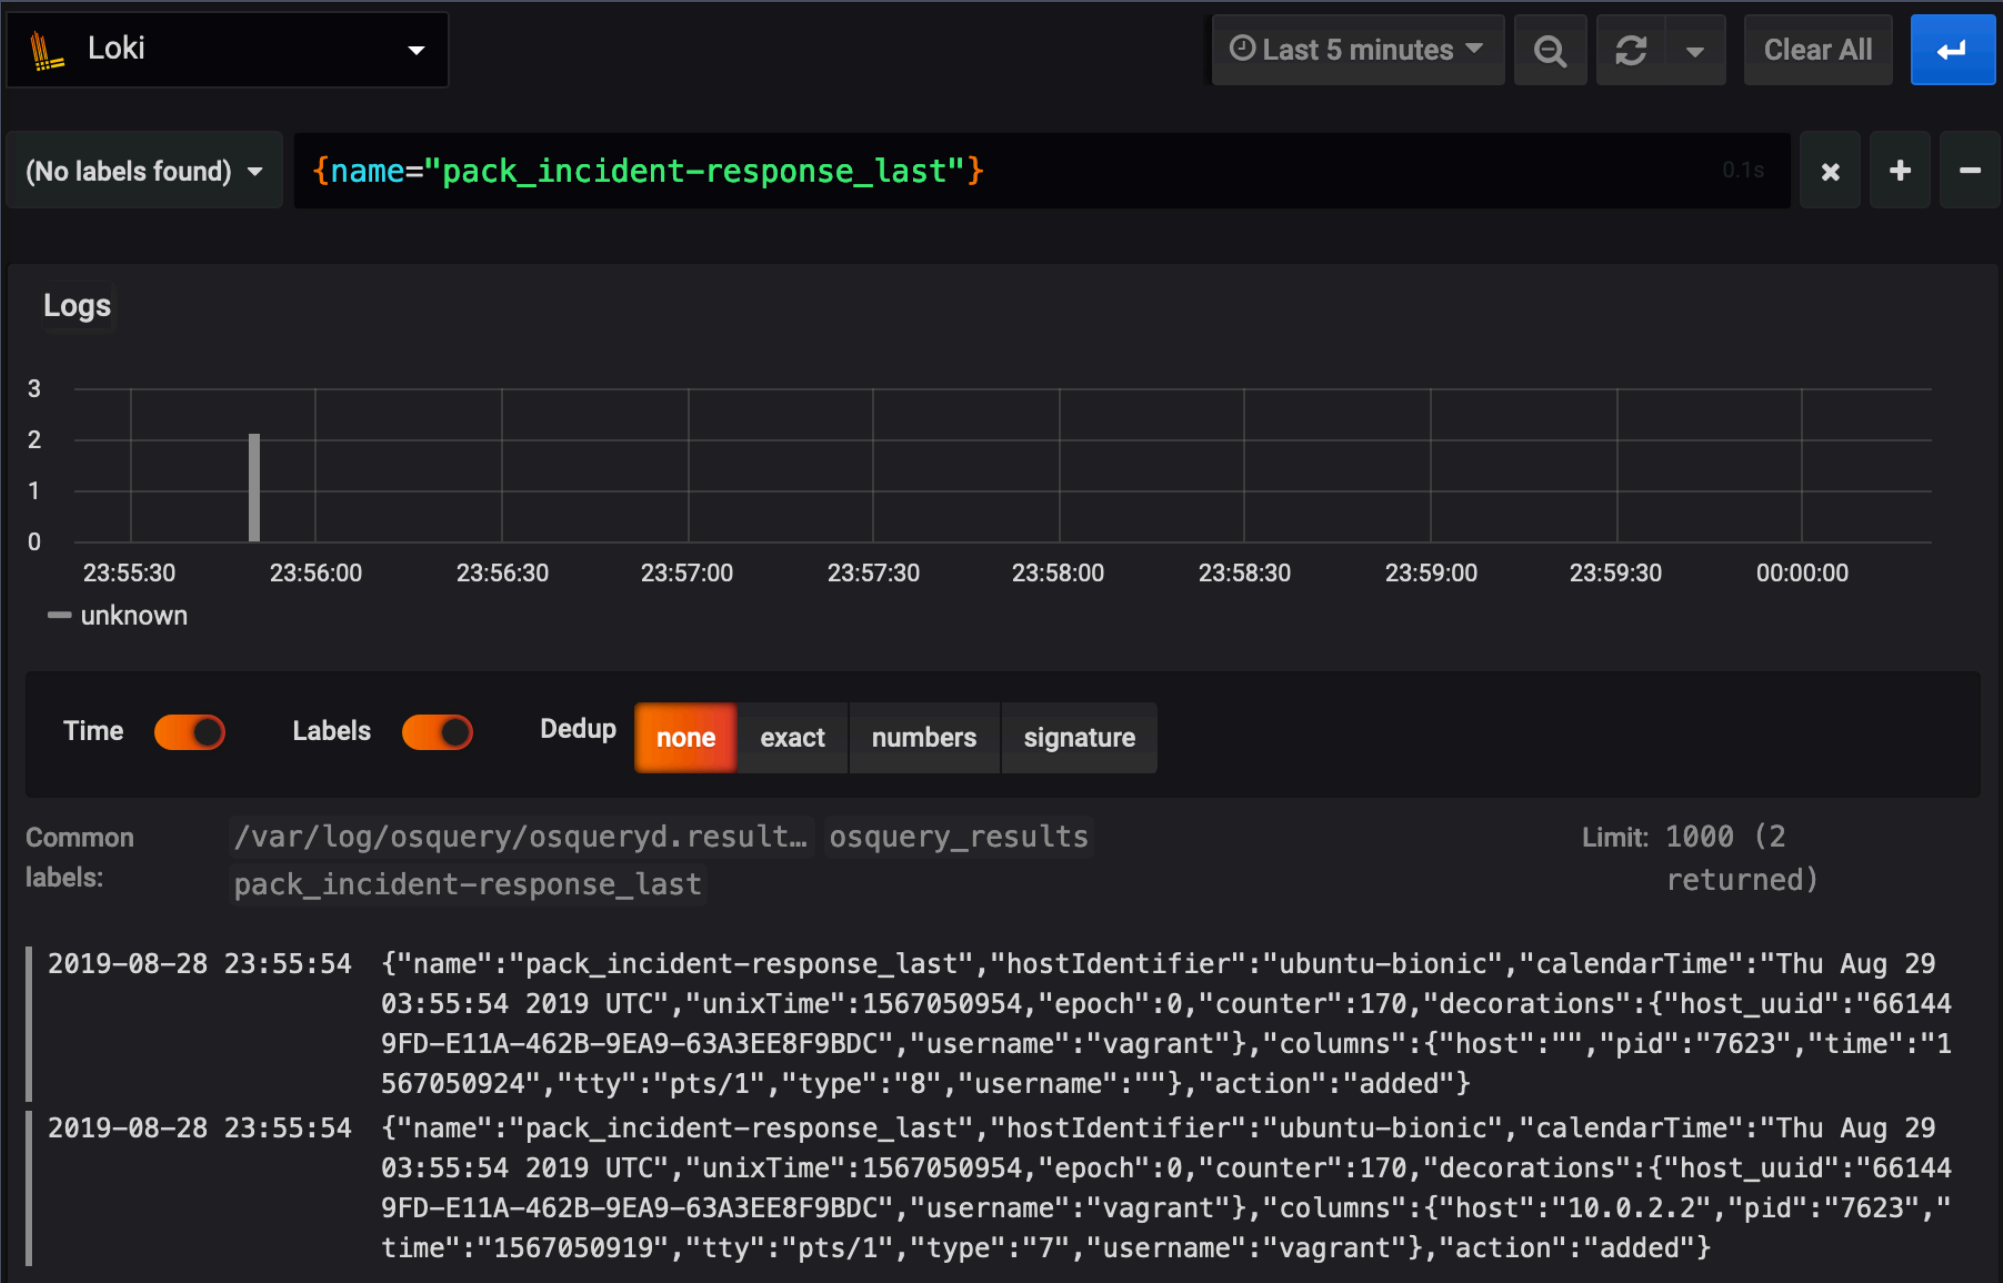
WHAT'S LEFT

CHARTING & ALERTING

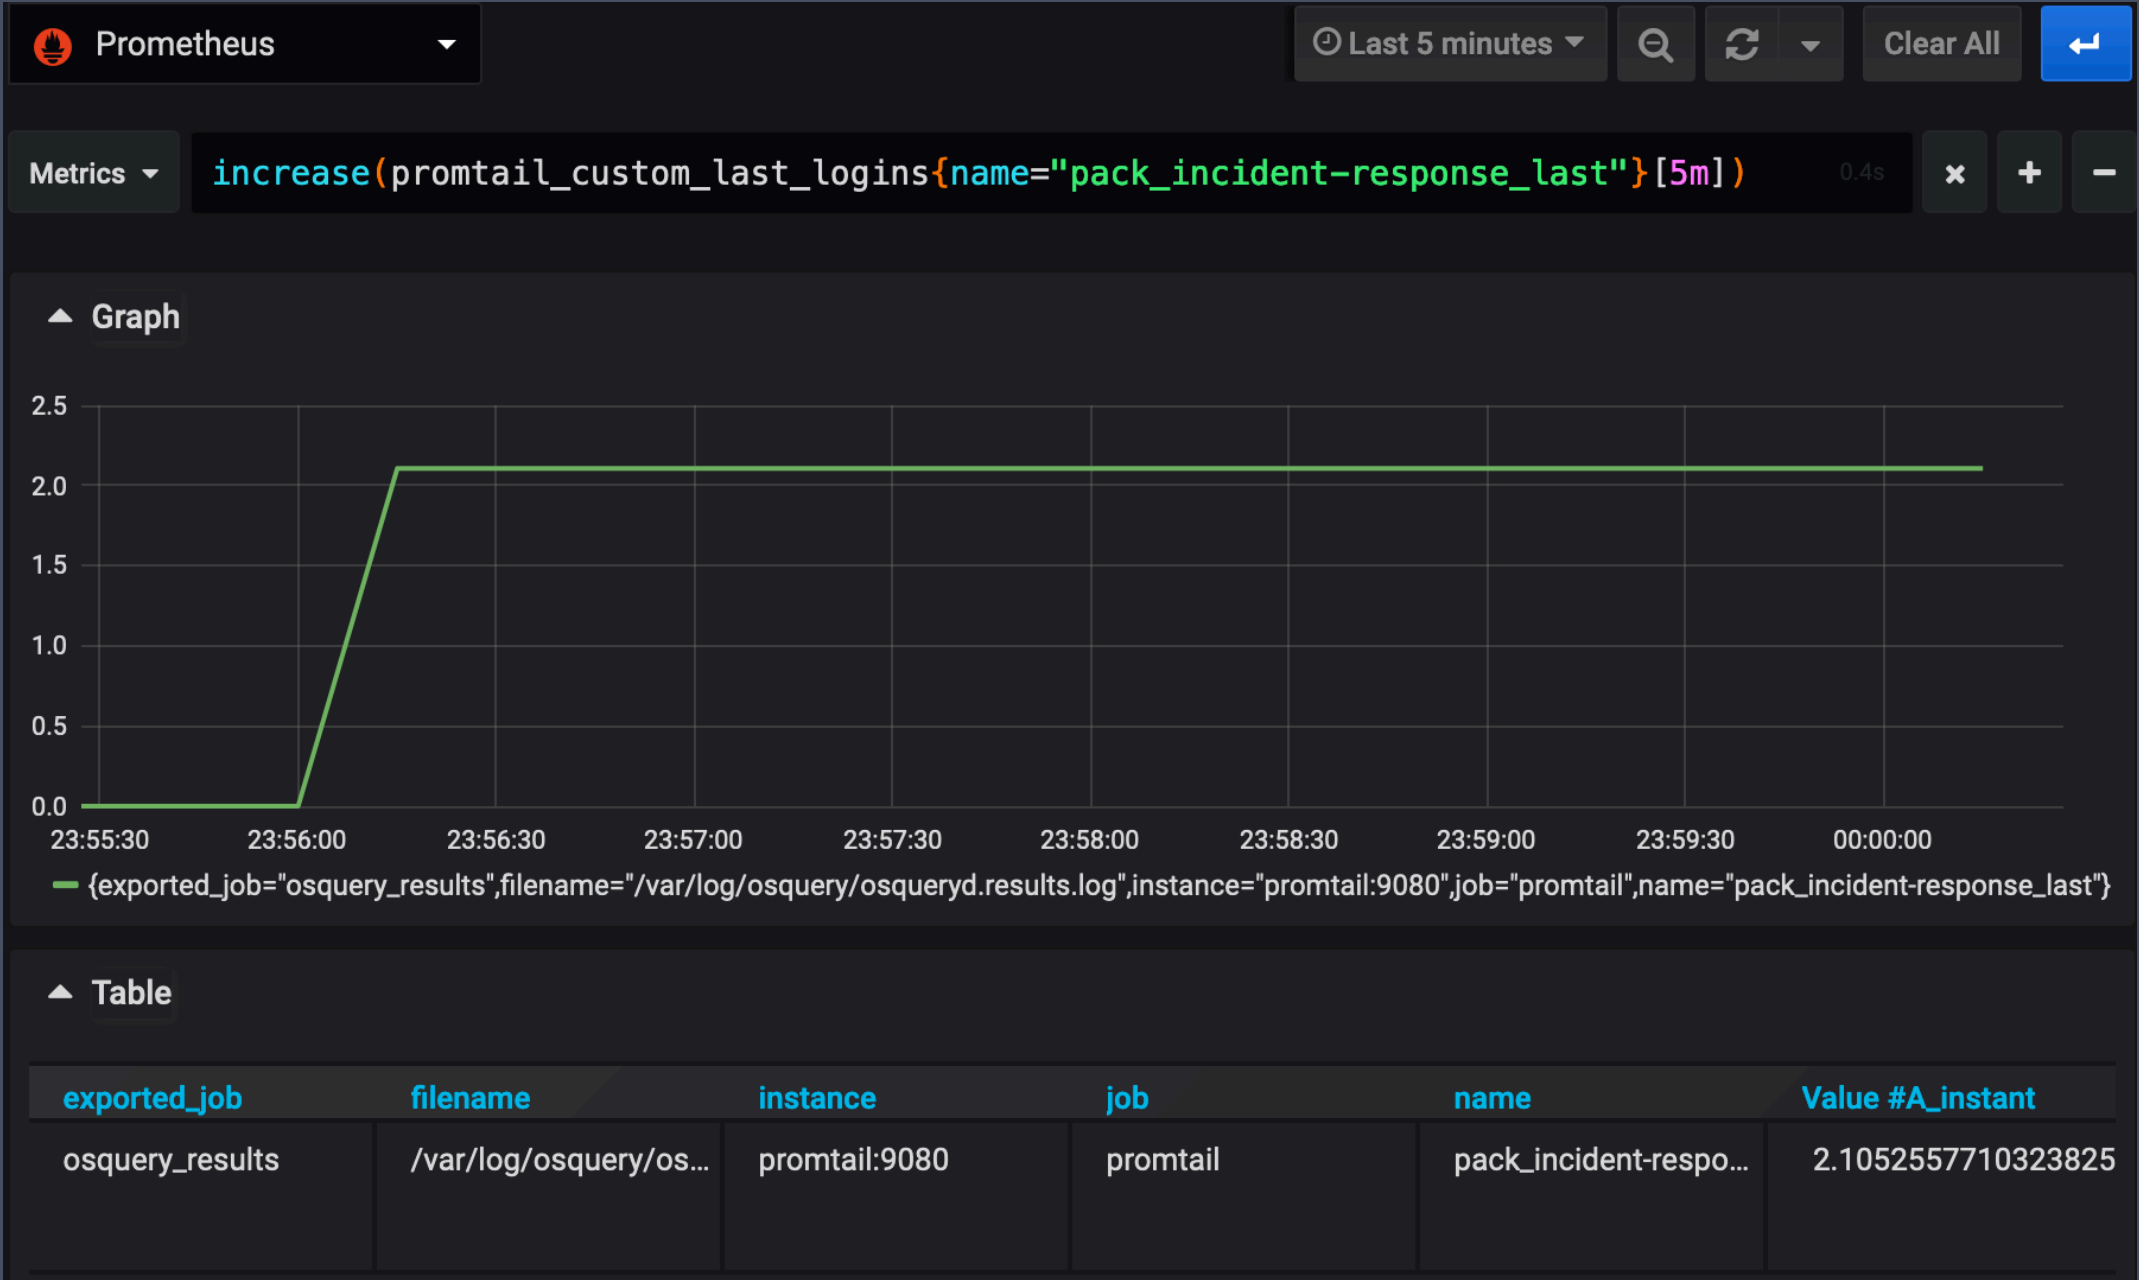
GRAFANA

SUPPORTS BOTH PROMTHEUS AND LOKI AS DATA SOURCES


GRAFANA - LOKI



GRAFANA – PROMETHEUS



GRAFANA - ALERTING



PUTTING IT ALL TOGETHER

38/38

TODO: diagram all the
components